


# Lecture 1 : Artificial Neural Networks

---

**Class:** ANN

**Date:** 07.58 AM,  Today

**Author:** Lasal Hettiarachchi

---

## Key learnings:

- ANN
- Forward Pass
- Backward Pass

## Real World Example

You go on shopping for a new Laptop. What are the factors you base your decision on? Assume all of these are binary variables

- The Price ( $X_p$ )
- Is it better than the Current Laptop. ( $X_b$ )
- Is the merchant reliable. ( $X_r$ )

Decision

$$W_1 X_p + W_2 X_b + W_3 X_r = y$$

- if  $y > 5 \rightarrow \text{yes}$
- if  $y \leq 5 \rightarrow \text{no}$

The weightes can be adjusted to make the decision

## Generalizing the Decision Statement

$$y = \begin{cases} 0, & x_P \cdot w_1 + x_b \cdot w_2 + x_r \cdot w_3 < t \\ 1, & x_P \cdot w_1 + x_b \cdot w_2 + x_r \cdot w_3 \geq t \end{cases}$$

- Its easier to compute with vectors


$$y = \begin{cases} 0, & \underline{x} \cdot \underline{w} < t \\ 1, & \underline{x} \cdot \underline{w} \geq t \end{cases}$$

$$y = \begin{cases} 0, & x \cdot w + b < 0 \\ 1, & x \cdot w + b \geq 0 \end{cases}$$

- We take the t and make it as a bias (bias is nothing but the negative of the threshold value)
- From this intuition we create an artificial neuron which takes in multiple features ( $x_P, x_b, x_r$  /  $X$ ) and bias b to output the vector y

What is an ANN?

- Collection of neuron is called an ANN
- When we stack neurons together , we get a layer



**Activation Function**










- Whats the purpose of activation functions?
  - Output of a perceptron is always binary, this is not good always
  - Linear outputs are useful but difficult to learn complex data
  - Non linear complex values are thus required
  - **To add non linearity**
- It is possible to transform any value between 0-1 using a linear activation function
- But if the output is linear, there isnt gonna be much change from layer to layer

$$\begin{aligned}
 a^{[1]} &= z^{[1]} = w^{[1]}x + b^{[1]} \\
 a^{[2]} &= z^{[2]} = w^{[2]}a^{[1]} + b^{[2]} \\
 a^{[2]} &= w^{[2]}(w^{[1]}x + b^{[1]}) + b^{[2]} \\
 &= \underbrace{(w^{[2]}w^{[1]})}_w x + \underbrace{(w^{[2]}b^{[1]} + b^{[2]})}_{b'} \\
 &= w'x + b'
 \end{aligned}$$

- This will just output a linear function of the input (**All layers of the neural network collapse into one** — with linear activation functions, no matter how many layers in the neural network, the last layer will be a linear function of the first layer)
- We cannot do back propagation aswell
- This helps the network to understand more complex datasets
- Helps normalize the outputs of the specific NN layer

Examples:

- Sigmoid (This is used in output layers mostly to get values between 0 and 1)
- tanh (This is better than sigmoid since the mean is 0, this centers the data to have 0 mean)

Name	Plot	Equation	Derivative
Identity		$f(x) = x$	$f'(x) = 1$
Binary step		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{for } x \neq 0 \\ ? & \text{for } x = 0 \end{cases}$
Logistic (a.k.a Soft step)		$f(x) = \frac{1}{1 + e^{-x}}$	$f'(x) = f(x)(1 - f(x))$
Tanh		$f(x) = \tanh(x) = \frac{2}{1 + e^{-2x}} - 1$	$f'(x) = 1 - f(x)^2$
ArcTan		$f(x) = \tan^{-1}(x)$	$f'(x) = \frac{1}{x^2 + 1}$
Rectified Linear Unit (ReLU)		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
Parameteric Rectified Linear Unit (PReLU) [2]		$f(x) = \begin{cases} \alpha x & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} \alpha & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
Exponential Linear Unit (ELU) [3]		$f(x) = \begin{cases} \alpha(e^x - 1) & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} f(x) + \alpha & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
SoftPlus		$f(x) = \log_e(1 + e^x)$	$f'(x) = \frac{1}{1 + e^{-x}}$

Issue with sigmoid and tanh: When x is very large or very small. The derivative is very small. This can slow down gradient descent

- Relu
- Leaky Relu

The advantage is that for most of the values other than 0, the derivative is not 0 thus gradient descent works well.

Properties of a good activation function

- Func should be continuous and differentiable everywhere
- Derivative should not saturate (tend toward 0) over its expected input range, since this may stall learning
- Derivatives should not explode. (This may cause numerical instability and pivot on a single layer)



### What, Why and Which?? Activation Functions

While studying neural network, we often come across the term—"Activation functions". What are these activation functions?  
<https://medium.com/@snaily16/what-why-and-which-activation-functions-b2bf748c0441>


Function	Advantage	Disadvantage
Sigmoid	The function is <b>differentiable</b> . That means, we can find the slope of the sigmoid curve at any two points.	<b>Vanishing gradient</b> — for very high or very low values of X, there is almost no change to the prediction, causing a vanishing gradient problem.
	<b>Output values bound</b> between 0 and 1, normalizing the output of each neuron.	Due to vanishing gradient problem, sigmoids have <b>slow convergence</b> .
		<b>Outputs not zero centered</b> .
		<b>Computationally expensive</b> .

Function	Advantage	Disadvantage
Tanh	<b>Zero centered</b> — making it easier to model inputs that have strongly negative, neutral, and strongly positive values.	It also suffers <b>vanishing gradient problem</b> and hence <b>slow convergence</b> .
	The function and its <b>derivative</b> both are <b>monotonic</b> .	
	Works better than sigmoid function	

Function	Advantage	Disadvantage
Relu	<b>Computationally efficient</b> — allows the network to converge very quickly	<b>The Dying ReLU problem</b> — when inputs approach zero, or are negative, the gradient of the function becomes zero, the network cannot perform back-propagation and cannot learn.
	<b>Non-linear</b> — although it looks like a linear function, ReLU has a derivative function and allows for back-propagation	

Function	Advantage	Disadvantage
Leaky Relu	<b>Prevents dying ReLU problem</b> — this variation of ReLU has a small positive slope in the negative area, so it does enable back-propagation, even for negative input values	<b>Results not consistent</b> — leaky ReLU does not provide consistent predictions for negative input values.
		During the front propagation if the learning rate is set very high it will overshoot killing the neuron.


## Perceptron architecture



### Feed forward NN

- Most ANNs are Single flow, uni directional
- No single loops or backtracking
  - CNN
  - MLP

### Forward pass



### Implementing the feed forward NN

Before we start training the data on the sigmoid neuron, We will build our model inside a class called *SigmoidNeuron*.



**Feedforward\_NeuralNetworks/FeedForwardNeuralNetwork.ipynb at master · Niranjankumar-**  
Build our neural networks from scratch. Contribute to Niranjankumar-c/Feedforward\_NeuralNetworks development by creati...  
[https://github.com/Niranjankumar-c/Feedforward\\_NeuralNetworks/blob/master/FeedForwardNetworks/FeedForwardNeural...](https://github.com/Niranjankumar-c/Feedforward_NeuralNetworks/blob/master/FeedForwardNetworks/FeedForwardNeural...)

```
class SigmoidNeuron:
    #initialization
    def __init__(self):
        self.w = None
        self.b = None
    #forward pass
    def perceptron(self, x):
        return np.dot(x, self.w.T) + self.b

    def sigmoid(self, x):
        return 1.0/(1.0 + np.exp(-x))
    #updating the gradients using mean squared error loss
    def grad_w_mse(self, x, y):
        y_pred = self.sigmoid(self.perceptron(x))
        return (y_pred - y) * y_pred * (1 - y_pred) * x

    def grad_b_mse(self, x, y):
        y_pred = self.sigmoid(self.perceptron(x))
        return (y_pred - y) * y_pred * (1 - y_pred)

    #updating the gradients using cross entropy loss
    def grad_w_ce(self, x, y):
        y_pred = self.sigmoid(self.perceptron(x))
        if y == 0:
            return y_pred * x
        elif y == 1:
            return -1 * (1 - y_pred) * x
        else:
            raise ValueError("y should be 0 or 1")

    def grad_b_ce(self, x, y):
        y_pred = self.sigmoid(self.perceptron(x))
        if y == 0:
            return y_pred
        elif y == 1:
            return -1 * (1 - y_pred)
        else:
            raise ValueError("y should be 0 or 1")
```

```

#model fit method
def fit(self, X, Y, epochs=1, learning_rate=1, initialise=True, loss_fn="mse",
display_loss=False):

    # initialise w, b
    if initialise:
        self.w = np.random.randn(1, X.shape[1])
        self.b = 0

    if display_loss:
        loss = {}

    for i in tqdm_notebook(range(epochs), total=epochs, unit="epoch"):
        dw = 0
        db = 0
        for x, y in zip(X, Y):
            if loss_fn == "mse":
                dw += self.grad_w_mse(x, y)
                db += self.grad_b_mse(x, y)
            elif loss_fn == "ce":
                dw += self.grad_w_ce(x, y)
                db += self.grad_b_ce(x, y)

        m = X.shape[1]
        self.w -= learning_rate * dw/m
        self.b -= learning_rate * db/m

        if display_loss:
            Y_pred = self.sigmoid(self.perceptron(X))
            if loss_fn == "mse":
                loss[i] = mean_squared_error(Y, Y_pred)
            elif loss_fn == "ce":
                loss[i] = log_loss(Y, Y_pred)

    if display_loss:
        plt.plot(loss.values())
        plt.xlabel('Epochs')
        if loss_fn == "mse":
            plt.ylabel('Mean Squared Error')
        elif loss_fn == "ce":
            plt.ylabel('Log Loss')
        plt.show()

def predict(self, X):
    Y_pred = []
    for x in X:
        y_pred = self.sigmoid(self.perceptron(x))
        Y_pred.append(y_pred)
    return np.array(Y_pred)

```

- The weight values are initialized randomly and the bias with 0



```

if initialise:
    self.w = np.random.randn(1, X.shape[1])
    self.b = 0

```

- Mean squared error with respect to w and updating gradients using cross entropy loss

```

#updating the gradients using mean squared error loss
def grad_w_mse(self, x, y):
    y_pred = self.sigmoid(self.perceptron(x))
    return (y_pred - y) * y_pred * (1 - y_pred) * x

#updating the gradients using cross entropy loss
def grad_w_ce(self, x, y):
    y_pred = self.sigmoid(self.perceptron(x))
    if y == 0:
        return y_pred * x
    elif y == 1:
        return -1 * (1 - y_pred) * x
    else:
        raise ValueError("y should be 0 or 1")

```

- Mean squared error with respect to b and updating gradients using cross entropy loss

```

def grad_b_mse(self, x, y):
    y_pred = self.sigmoid(self.perceptron(x))
    return (y_pred - y) * y_pred * (1 - y_pred)

def grad_b_ce(self, x, y):
    y_pred = self.sigmoid(self.perceptron(x))
    if y == 0:
        return y_pred
    elif y == 1:
        return -1 * (1 - y_pred)
    else:
        raise ValueError("y should be 0 or 1")

```

## Loss function (Cost)

- Parameter that measures how well the NN is doing
- Goal is to find w , b that minimizes the loss
- No of different loss functions are available
  - MSE
  - RSME
  - Hinge Loss
- This is the different between the expected label in the training set and the predicted value

Loss with respect to w can be calculated by when the function is sigmoid,

$$\begin{aligned}
\nabla w &= \frac{\partial}{\partial w} \left[ \frac{1}{2} * (f(x) - y)^2 \right] \\
&= \frac{1}{2} * [2 * (f(x) - y) * \frac{\partial}{\partial w} (f(x) - y)] \\
&= (f(x) - y) * \frac{\partial}{\partial w} (f(x)) \\
&= (f(x) - y) * \frac{\partial}{\partial w} \left( \frac{1}{1 + e^{-(wx+b)}} \right) \\
&= (f(x) - y) * f(x) * (1 - f(x)) * x
\end{aligned}
\quad \left| \quad
\begin{aligned}
&\frac{\partial}{\partial w} \left( \frac{1}{1 + e^{-(wx+b)}} \right) \\
&= \frac{-1}{(1 + e^{-(wx+b)})^2} \frac{\partial}{\partial w} (e^{-(wx+b)}) \\
&= \frac{-1}{(1 + e^{-(wx+b)})^2} * (e^{-(wx+b)}) \frac{\partial}{\partial w} (-(wx+b)) \\
&= \frac{-1}{(1 + e^{-(wx+b)})} * \frac{e^{-(wx+b)}}{(1 + e^{-(wx+b)})} * (-x) \\
&= \frac{1}{(1 + e^{-(wx+b)})} * \frac{e^{-(wx+b)}}{(1 + e^{-(wx+b)})} * (x) \\
&= f(x) * (1 - f(x)) * x
\end{aligned}$$

- $W[1]$  and  $b[1]$  are weight and bias vectors of hidden layer 1 while ,  $W[2]$  and  $b[2]$  are weight and bias vectors of hidden layer 2.
- We use gradient decent to update the Vectors

## Gradient descent for neural networks

Parameters:  $W^{[1]}, b^{[1]}, W^{[2]}, b^{[2]}$   
 $(n^{[1]}, n^{[2]})$   $(n^{[2]}, n^{[3]})$   $(n^{[3]}, 1)$   $n_x = n^{[0]}, n^{[1]}, n^{[2]} = 1$

Cost function:  $J(W^{[1]}, b^{[1]}, W^{[2]}, b^{[2]}) = \frac{1}{m} \sum_{i=1}^m \ell(\hat{y}_i, y_i)$   
 $\uparrow a^{[2]}$

Gradient descent:

Repeat {


Compute predicts  $(\hat{y}^{(i)}, i=1, \dots, m)$

$dW^{[1]} = \frac{\partial J}{\partial W^{[1]}}$ ,  $db^{[1]} = \frac{\partial J}{\partial b^{[1]}}$ , ...

$W^{[1]} = W^{[1]} - \alpha dW^{[1]}$

$b^{[1]} = b^{[1]} - \alpha db^{[1]}$

# Logistic regression derivatives



- $\mathcal{L}$  is the loss. we need to minimize loss WRT  $w_1, w_2, b$

## Process of an FF NN

1. Randomly initialize weights
2. Implement forward propagation
3. Compute error total  $E_{total}$
4. Implement back propagation  $\frac{dE_{total}}{dw_j}$  (Etotal) /  $\frac{dE_{total}}{dw_j}$  (Wjk)
5. Use gradient decent or any other optimization technique to update the weights
6. Repeat the process over multiple iterations or error converges